




FiLM: each hidden layer's activation is rescaled by $\gamma_i(\alpha)$ and shifted by $\beta_i(\alpha)$ before the nonlinearity — the network's function of r changes continuously with α , the same way a parametric dispersion model is evaluated at any wavelength without refitting.