




Same shared backbone for Experiment 2 (soft label $\hat{\rho} \approx \rho_{\text{vox}}$) and Experiment 3 (hard label $\hat{\rho} \in \{0, 1\}$)
— only the training target and label binarisation change; architecture and FiLM mechanism are identical.